

Statistical Methods for Discrete Response, Time Series, and Panel Data (W271): Lab 3

K.C. Tobin, Weixing Sun, Winston Lin

August 6, 2017

Question 1: EDA

During your EDA, you notice that your data exhibits both seasonality (different months have different heights) AND that there is a clear linear trend. How many order of non-seasonal and seasonal differencing would it take to make this time-series stationary in the mean? Why?

Behavior of ACF and PACF for ARMA models (from SS2016, p.71): - AR(p): ACF tails off, PACF cuts off after lag p - MA(q): ACF cuts off after lag q, PACF tails off - ARMA(p,q): both ACF and PACF tail off

Here are the steps to building an ARIMA model: 1. Conduct an EDA to determine if you need to transform the data in order to make it stationary. 2. Transform the data if needed. 3. Estimate several Arima(p,d,q) models. Remember, you set the value of d in the first step! So really, you are trying to find the appropriate values of p and q. 4. Evaluate the residuals of models with the lowest AIC/BIC values and simpler models. Select the model where the residuals resemble white noise. 5. If you still have some candidate models remaining, then conduct an out of sample test and select the model with the lowest forecasting error. 6. Answer your question / generate forecasts!

```
### 1. EDA for unemployment rate
unemp = read.csv("UNRATENSA.csv")
cbind(head(unemp),tail(unemp))
```


##	DATE	UNRATENSA	DATE	UNRATENSA
## 1	1948-01-01	4.0	2017-01-01	5.1
## 2	1948-02-01	4.7	2017-02-01	4.9
## 3	1948-03-01	4.5	2017-03-01	4.6
## 4	1948-04-01	4.0	2017-04-01	4.1
## 5	1948-05-01	3.4	2017-05-01	4.1
## 6	1948-06-01	3.9	2017-06-01	4.5

```
unemp.ts = ts(unemp$UNRATENSA, start = c(1948, 1), frequency=12)
```

```
# plot series
```



```
plot.ts(unemp.ts, main = 'Monthly unemployment rate, January 1948 - June 2017', ylab = 'Unemployment rate')
```

Monthly unemployment rate, January 1948 – June 2017



```
invisible(acf2(unemp.ts, 130)) # ACF: trend, PACF: seasonality
```

Series: unemp.ts

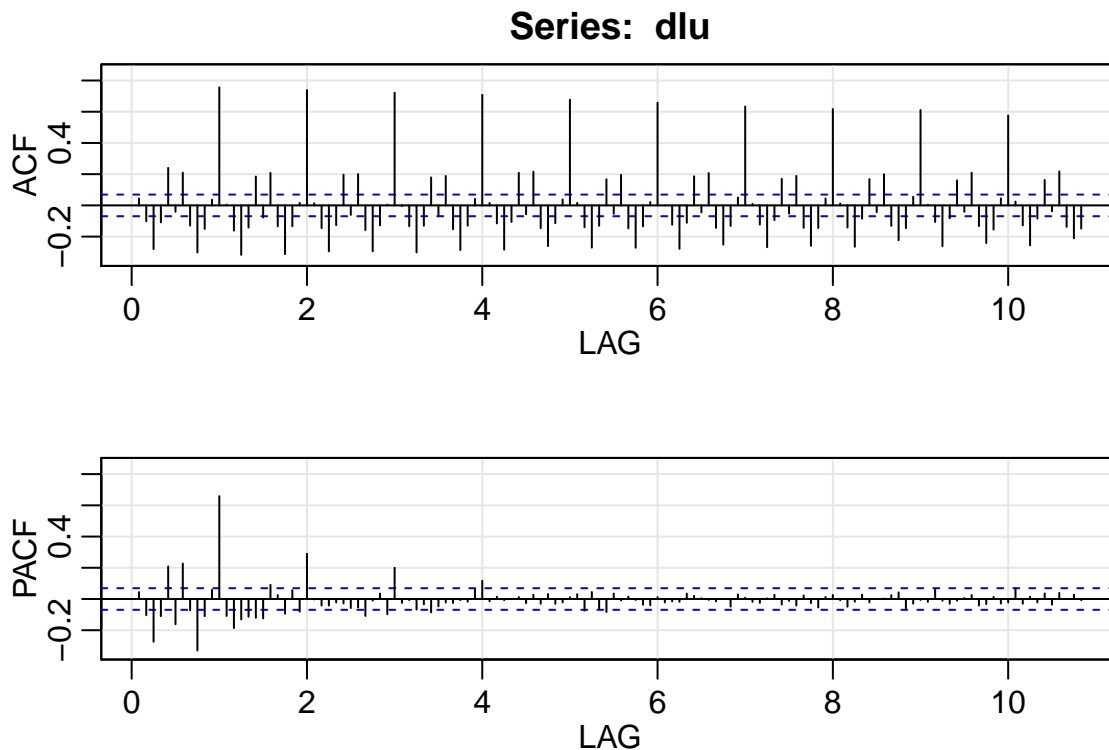


```
# test null hypothesis of non-stationarity (unit root)
adf.test(unemp.ts) # p = 0.33 > 0.05, series is non-stationary
```

```
##
```

```
## Augmented Dickey-Fuller Test
##
## data: unemp.ts
## Dickey-Fuller = -2.5911, Lag order = 9, p-value = 0.3281
## alternative hypothesis: stationary
### 2. transform data to stationary series
# log transform to stabilize variance
lu = log(unemp.ts)

# diff by 1 lag to remove monthly trend
dlu = diff(lu)
invisible(acf2(dlu, 130))
```





```
# ACF: seasonal trend tails off --> seasonal MA
# PACF: cutoff after 3 seasonal lags --> seasonal AR, max order 3

# test null hypothesis of non-stationarity (unit root)
adf.test(dlu)

## Warning in adf.test(dlu): p-value smaller than printed p-value
##
## Augmented Dickey-Fuller Test
##
## data: dlu
## Dickey-Fuller = -13.715, Lag order = 9, p-value = 0.01
## alternative hypothesis: stationary
# p = 0.01 < 0.05 --> series is stationary, but seasonal MA from PACF
# diff by 12 lags to remove seasonality
```

```
ddlu = diff(dlu, lag=12)
invisible(acf2(ddlu, 130))
```



```
# ACF: cutoff after 1 seasonal lag --> seasonal MA, max order 1
# ACF: cutoff after 3 lags --> MA, max order 3
# PACF: cutoff after 3 seasonal lags --> seasonal
```

```
# test null hypothesis of non-stationarity (unit root)
adf.test(ddlu)
```

```
## Warning in adf.test(ddlu): p-value smaller than printed p-value
```

```
##
```

```
## Augmented Dickey-Fuller Test
```

```
##
```

```
## data: ddlu
```

```
## Dickey-Fuller = -9.7378, Lag order = 9, p-value = 0.01
```


```
## alternative hypothesis: stationary
```

```
# p = 0.01 < 0.05, series is stationary
```

```
# plot transformed data
```

```
plot.ts(cbind(unemp.ts,dlu,ddlu), main='Monthly unemployment rate, January 1948 - June 2017')
```

Monthly unemployment rate, January 1948 – June 2017



start w/ d=1, D=1 from ADF tests for stationarity, max of 3 seasonal and 3 non-seasonal lags from ACF

The variance of the series appears to decrease over time, so we take the log of the series to stabilize the variance. The ACF decays slowly, so we take the difference with lag 1 (1 month) to remove the trend. The resulting ACF shows seasonal autocorrelation every 12 months that decays gradually. Now, we take the difference with lag 12 (1 year) to remove the seasonal trend. The ACF now cuts off after lag 12, and the PACF decays slowly. This hints that a seasonal MA(1) model might be a good place to start fitting. So to remove trend and seasonality, we would perform non-seasonal differencing of order 1 and seasonal differencing of order 1, with a 12-month seasonal period.

```
### 1. EDA for automotive sales
auto = read.csv("TOTALNSA.csv")
cbind(head(auto), tail(auto))
```


##	DATE	TOTALNSA	DATE	TOTALNSA
## 1	1976-01-01	885.2	2017-01-01	1164.3
## 2	1976-02-01	994.7	2017-02-01	1352.1
## 3	1976-03-01	1243.6	2017-03-01	1582.7
## 4	1976-04-01	1191.2	2017-04-01	1449.7
## 5	1976-05-01	1203.2	2017-05-01	1544.1
## 6	1976-06-01	1254.7	2017-06-01	1500.6

```
auto.ts = ts(auto$TOTALNSA, start = c(1976, 1), frequency=12)
```

```
# plot raw data
```



```
plot.ts(auto.ts, main = 'Monthly automotive sales, January 1976 - June 2017', ylab = 'Automotive sales')
```

Monthly automotive sales, January 1976 – June 2017



```
invisible(acf2(auto.ts, 130)) # ACF: trend and seasonality
```

Series: auto.ts

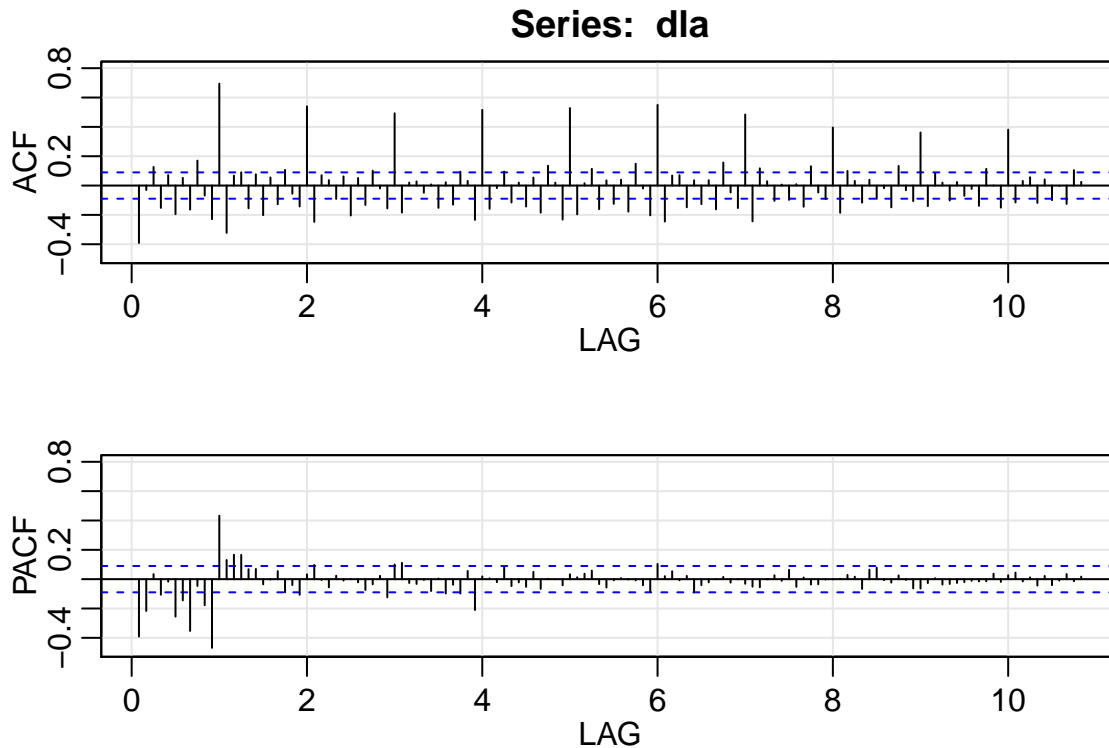


```
# test null hypothesis of non-stationarity (unit root)  
adf.test(auto.ts) # p = 0.04 < 0.05, series is stationary, no transform needed
```

```
##
```

```
## Augmented Dickey-Fuller Test
##
## data: auto.ts
## Dickey-Fuller = -3.5662, Lag order = 7, p-value = 0.03595
## alternative hypothesis: stationary
## 2. transform data to stationary series


# diff by 1 lag to remove monthly trend
dla = diff(auto.ts)
invisible(acf2(dla, 130))
```



```
# test null hypothesis of non-stationarity (unit root)
adf.test(dla) # p = 0.01 < 0.05, series is stationary
```

```
## Warning in adf.test(dla): p-value smaller than printed p-value
```

```
##
## Augmented Dickey-Fuller Test
##
## data: dla
## Dickey-Fuller = -16.651, Lag order = 7, p-value = 0.01
## alternative hypothesis: stationary
# diff by 12 lags to remove seasonal trend
ddla = diff(dla, lag=12)
invisible(acf2(ddla, 250))
```




```
# test null hypothesis of non-stationarity (unit root)
adf.test(ddla) # p = 0.01 < 0.05, series is stationary

## Warning in adf.test(ddla): p-value smaller than printed p-value
##
## Augmented Dickey-Fuller Test
##
## data: ddla
## Dickey-Fuller = -11.281, Lag order = 7, p-value = 0.01
## alternative hypothesis: stationary

# plot transformed data
plot.ts(cbind(auto.ts,dla,ddla), main='Monthly automotive sales, January 1976 - June 2017')
```


Monthly automotive sales, January 1976 – June 2017



The variance of the series appears to be relatively constant so no need to apply a transformation. The ACF decays slowly, so we take the difference with lag 1 (1 month) to remove the trend. The resulting ACF shows seasonal autocorrelation every 12 months that decays gradually. Thus, we take the difference with lag 12 (1 year) to remove the seasonal trend. The ACF now cuts off after lag 12, and the PACF decays slowly. So to remove trend and seasonality, we would perform non-seasonal differencing of order 1 and seasonal differencing of order 1, with a 12-month seasonal period.

Question 2: SARIMA

It is Dec. 31, 2016, and you work for a non-partisan think tank focusing on the state of the U.S. economy. You are interested in forecasting the unemployment rate through 2017 (and then 2020) to use it as a benchmark against the incoming administration's economic performance. Use the dataset UNRATENSA.csv and answer the following:

- Build a SARIMA model using the unemployment data and produce a 1-year forecast and then a 4-year forecast. Because it is Dec. 31, 2016, leave out 2016 as your test data.

```
### 3.a.fit seasonal models and select final candidate models with lowest AIC/BIC
```

```
# split data into training and test sets
```

```
unemp.train = window(unemp.ts, end=c(2015,12), frequency=12)
```

```
unemp.test = window(unemp.ts, start=c(2016,1), frequency=12)
```

```
# try models w/ max 3 seasonal AR, MA lags based on ACF, PACF
```

```
# start w/ D=1 from EDA and transforms
```

```
for (P in 0:3) {
```

```
  for(Q in 0:3) {
```

```
    fit = Arima(unemp.train, order=c(0,1,0),
```

```
               seasonal=list(order=c(P,1,Q)), method='ML')
```

```

        print(c(P,Q,fit$aic))
    }
}

```

```

## [1] 0.0000 0.0000 427.2262
## [1] 0.00000 1.00000 47.12622
## [1] 0.00000 2.00000 49.10322
## [1] 0.00000 3.00000 51.10067
## [1] 1.0 0.0 210.4
## [1] 1.00000 1.00000 49.10312
## [1] 1.00000 2.00000 51.12536
## [1] 1.00000 3.00000 52.14965
## [1] 2.0000 0.0000 138.9131
## [1] 2.00000 1.00000 51.10083
## [1] 2.00000 2.00000 52.21086
## [1] 2.00000 3.00000 42.59918
## [1] 3.0000 0.0000 85.7188
## [1] 3.00000 1.00000 52.90777
## [1] 3.00000 2.00000 53.35645
## [1] 3.00000 3.00000 44.58169

```

```

# candidate seasonal models (P,D,Q):
# (0,1,1) AIC = 47.1 --> low AIC, largest decrease in AIC, most parsimonious model
# (2,1,3) AIC = 42.6 --> lowest AIC, not much improvement in AIC, more complex model

```

Based on the AIC outputs it appears that a seasonal order of (0,1,1) produces relatively good results and makes for a parsimonious choice to move forward with.

3.b. fit non-seasonal models (given seasonal models) and select final candidate models with lowest AIC

```

# try models w/ max of 3 non-seasonal AR, MA lags based on ACF, PACF
for (p in 0:3) {
  for(q in 0:3) {
    try(fit <- Arima(unemp.train, order=c(p,1,q),
                    seasonal=list(order=c(0,1,1)), method='ML'))
    print(c(p,q,fit$aic))
  }
}

```

```

## [1] 0.00000 0.00000 47.12622
## [1] 0.00000 1.00000 32.70797
## [1] 0.000000 2.000000 1.640301
## [1] 0.000000 3.000000 -2.766081
## [1] 1.00000 0.00000 25.49529
## [1] 1.00000 1.00000 -12.69638
## [1] 1.00000 2.00000 -18.01948
## [1] 1.00000 3.00000 -16.24685
## [1] 2.00000 0.00000 -10.45144
## [1] 2.00000 1.00000 -18.34985
## [1] 2.00000 2.00000 -16.39718
## [1] 2.00000 3.00000 -14.39471
## [1] 3.00000 0.00000 -15.31482
## [1] 3.00000 1.00000 -16.39161
## [1] 3.0000 2.0000 -14.3607
## [1] 3.00000 3.00000 -15.81192

```

```
# candidate non-seasonal models (p,d,q):
# (0,1,0) AIC = 47.1 --> baseline model
# (0,1,2) AIC = 1.6 --> low AIC, largest decrease in AIC, most parsimonious model
# (1,1,1) AIC = -12.7 --> low AIC, large decrease in AIC
# (1,1,2) AIC = -18.0 --> 2nd lowest AIC, slight improvement
# (2,1,1) AIC = -18.3 --> lowest AIC, slight improvement
# (2,1,2) AIC = -16.4
```


Several AR/MA terms are close in nature so we will continue to evaluate them all below.

```
### 4. check candidate models for residuals ~ white noise
sarima_diag = function(train, p,d,q, P,D,Q) {
  m = Arima(train, order=c(p,d,q),
            seasonal=list(order=c(P,D,Q)), method='ML')


  # residual diagnostics
  par(mfcol=c(2,2))
  hist(m$residuals)
  qqnorm(m$residuals)
  acf(m$residuals, 130)
  pacf(m$residuals, 130)
}

# diagnostics for candidate models
sarima_diag(unemp.train, 0,1,0, 0,1,1) # residuals are serially correlated, normally distributed
```


Histogram of m\$residuals




Series m\$residuals



Normal Q-Q Plot




Series m\$residuals




```
sarima_diag(unemp.train, 0,1,2, 0,1,1) # residuals are serially uncorrelated, normally distributed
```


Histogram of m\$residuals




Series m\$residuals






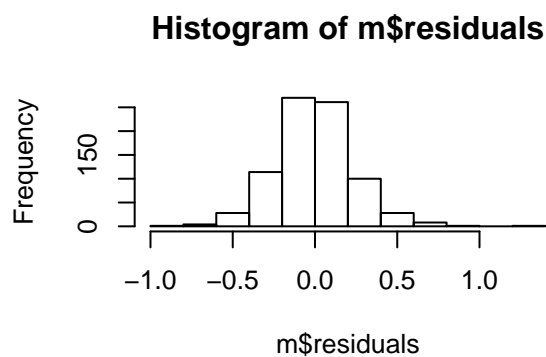
Normal Q-Q Plot







Series m\$residuals



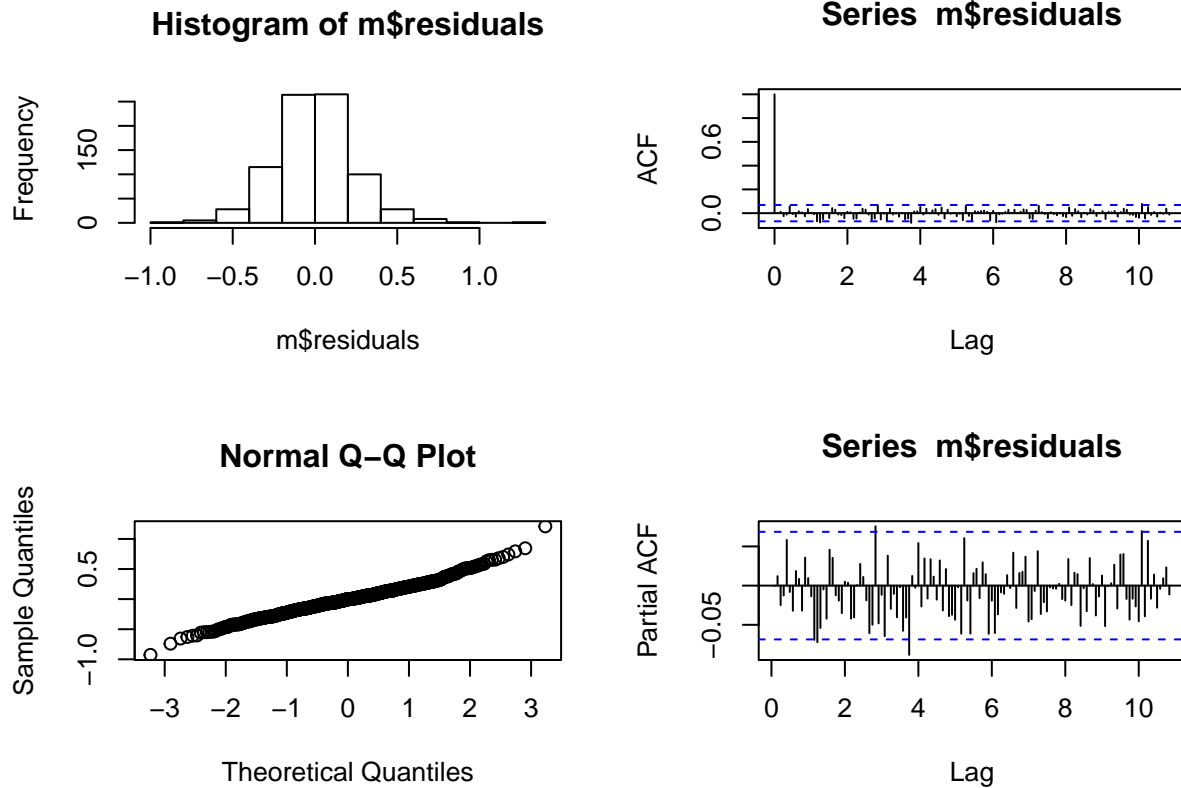
```
sarima_diag(unemp.train, 2,1,1, 0,1,1) # residuals are serially uncorrelated, normally distributed
```



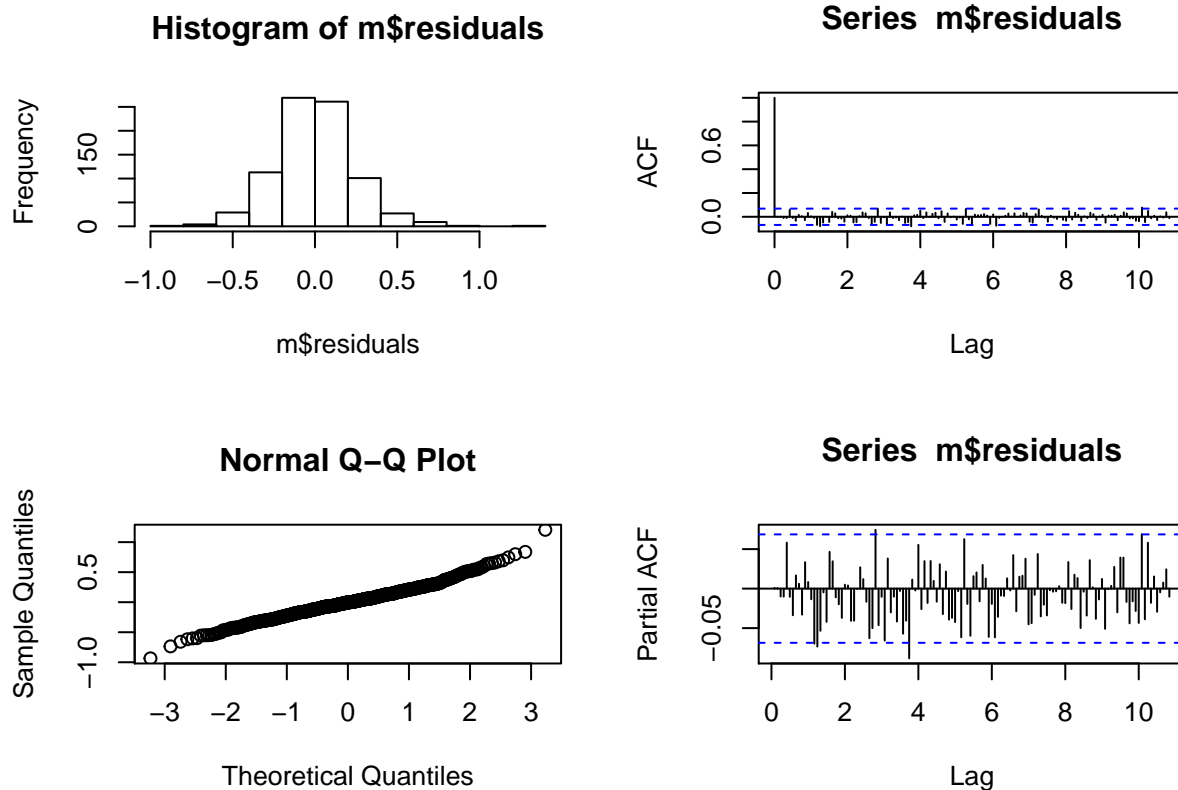
```
sarima_diag(unemp.train, 1,1,1, 0,1,1) # residuals are serially uncorrelated, normally distributed
```



```
sarima_diag(unemp.train, 1,1,2, 0,1,1) # residuals are serially uncorrelated, normally distributed
```



```
sarima_diag(unemp.train, 2,1,2, 0,1,1) # residuals are serially uncorrelated, normally distributed
```



All models appear to produce residuals that resemble white noise and lack correlation with the exception of that (0,1,0) model. We will now evaluate the various models in terms of their forecasting power.

5. check candidate models for out of sample accuracy

```
# candidate models
m0 = Arima(unemp.train, order=c(0,1,0), seasonal=list(order=c(0,1,1)), method='ML')
m1 = Arima(unemp.train, order=c(0,1,2), seasonal=list(order=c(0,1,1)), method='ML')
m2 = Arima(unemp.train, order=c(1,1,1), seasonal=list(order=c(0,1,1)), method='ML')
m3 = Arima(unemp.train, order=c(1,1,2), seasonal=list(order=c(0,1,1)), method='ML')
m4 = Arima(unemp.train, order=c(2,1,1), seasonal=list(order=c(0,1,1)), method='ML')
m5 = Arima(unemp.train, order=c(2,1,2), seasonal=list(order=c(0,1,1)), method='ML')
```

```
# forecasts for test period
for0 = forecast(m0, h=length(unemp.test))
for1 = forecast(m1, h=length(unemp.test))
for2 = forecast(m2, h=length(unemp.test))
for3 = forecast(m3, h=length(unemp.test))
for4 = forecast(m4, h=length(unemp.test))
for5 = forecast(m5, h=length(unemp.test))
```

```
# check accuracy against test set
accuracy(for0, unemp.test) # RMSE = 0.152
```

```
##              ME      RMSE      MAE      MPE      MAPE
## Training set -0.006361257 0.2447366 0.1825672 -0.08723024 3.367421
## Test set    -0.018362544 0.1518747 0.1272645 -0.40584864 2.699796
##              MASE      ACF1 Theil's U
## Training set 0.2063602 0.1685593      NA
```

```
## Test set      0.1438503 0.4257728 0.4583755
```

```
accuracy(for1, unemp.test) # RMSE = 0.153
```

```
##              ME      RMSE      MAE      MPE      MAPE
## Training set -0.004965367 0.2374515 0.1793519 -0.05758716 3.329680
## Test set     0.026112182 0.1530490 0.1196016  0.55779723 2.523776
##              MASE      ACF1 Theil's U
## Training set 0.2027259 0.01851124      NA
## Test set     0.1351887 0.39501668  0.454131
```

```
accuracy(for2, unemp.test) # RMSE = 0.155
```

```
##              ME      RMSE      MAE      MPE      MAPE
## Training set -0.003388788 0.2354166 0.1791475 -0.004004457 3.334639
## Test set     0.034894521 0.1547726 0.1217408  0.750543864 2.567508
##              MASE      ACF1 Theil's U
## Training set 0.2024949 -0.05554858      NA
## Test set     0.1376066 0.38914276 0.4598686
```

```
accuracy(for3, unemp.test) # RMSE = 0.151 --> best fit, use SARIMA(1,1,2)(0,1,1) for final forecast
```

```
##              ME      RMSE      MAE      MPE      MAPE
## Training set -0.003534511 0.2343297 0.1778710 -0.01030263 3.320626
## Test set     0.015050078 0.1506237 0.1193972  0.32696024 2.518034
##              MASE      ACF1 Theil's U
## Training set 0.2010520 -0.0004805127      NA
## Test set     0.1349576 0.3868569700 0.4453878
```

```
accuracy(for4, unemp.test) # RMSE = 0.152
```

```
##              ME      RMSE      MAE      MPE      MAPE
## Training set -0.003512207 0.2342802 0.1777658 -0.009414411 3.318480
## Test set     0.017207308 0.1510337 0.1194946  0.372832772 2.520191
##              MASE      ACF1 Theil's U
## Training set 0.2009331 0.001441396      NA
## Test set     0.1350677 0.387328132  0.446624
```

```
accuracy(for5, unemp.test) # RMSE = 0.151
```


```
##              ME      RMSE      MAE      MPE      MAPE
## Training set -0.00349946 0.2342739 0.1777561 -0.008939798 3.317892
## Test set     0.01817224 0.1512012 0.1195267  0.393400264 2.520929
##              MASE      ACF1 Theil's U
## Training set 0.2009221 0.0009387734      NA
## Test set     0.1351040 0.3873410157 0.4471562
```

Based on forecasting power Sarima (1,1,2)(0,1,1) performs slightly better and will be our choice for modeling unemployment rates.

- How well does your model predict the unemployment rate up until June 2017?

The SARIMA(1,1,2) (0,1,1) model predicts unemployment rate well, with a RMSE of 0.151 on the test data set. Forecast seems to reflect the trend and seasonality and does not deviate significantly from the actual values in the forecast horizon.


```
### 6.a. forecast through June 2017
# final forecast with model 3: SARIMA(1,1,2) (0,1,1) and CIs
for2y = sarima.for(unemp.train, length(unemp.test), 1,1,2, 0,1,1, 12)
lines(unemp.test, col='black')
```

- What does the unemployment rate look like at the end of 2020? How credible is this estimate?

The model predicts an unemployment rate of 2.53% at the end of 2020. The point estimate appears somewhat credible based on the current downward trend but would still reflect an all time low. In addition the 95% confidence interval widens and contains 0 and negative values as well as record highs in December 2020.

```
### 6.a. forecast through December 2020
# final forecast with model 3: SARIMA(1,1,2) (0,1,1) and CIs
for5y = sarima.for(unemp.train, 60, 1,1,2, 0,1,1, 12)
lines(unemp.test, col='black')
```



```
tail(for5y)
```

```
## $pred
##           Jan      Feb      Mar      Apr      May      Jun      Jul
## 2016 5.526246 5.357901 5.115507 4.533402 4.686756 5.055630 5.146009
## 2017 5.079234 4.908464 4.664247 4.080770 4.233092 4.601191 4.690987
## 2018 4.622763 4.451914 4.207637 3.624115 3.776404 4.144477 4.234254
## 2019 4.165983 3.995132 3.750852 3.167329 3.319617 3.687689 3.777465
## 2020 3.709193 3.538341 3.294062 2.710539 2.862826 3.230898 3.320675
##           Aug      Sep      Oct      Nov      Dec
## 2016 4.856990 4.502899 4.383385 4.337381 4.360586
## 2017 4.401529 4.047109 3.927347 3.881156 3.904221
## 2018 3.944782 3.590351 3.470581 3.424384 3.447444
## 2019 3.487993 3.133561 3.013791 2.967594 2.990654
## 2020 3.031202 2.676771 2.557000 2.510803 2.533863
##
## $se
##           Jan      Feb      Mar      Apr      May      Jun      Jul
## 2016 0.2362185 0.3531598 0.4680527 0.5789404 0.6849151 0.7856838 0.8812959
## 2017 1.3821818 1.4693835 1.5555925 1.6401072 1.7225385 1.8027004 1.8805360
## 2018 2.3161690 2.3945204 2.4731011 2.5512692 2.6285951 2.7048032 2.7797268
## 2019 3.2118738 3.2898950 3.3684326 3.4469109 3.5249268 3.6022056 3.6785663
## 2020 4.1243983 4.2045009 4.2851985 4.3659668 4.4464315 4.5263308 4.6054864
##           Aug      Sep      Oct      Nov      Dec
## 2016 0.9719833 1.0580661 1.1398982 1.2178360 1.2922212
## 2017 1.9560681 2.0293674 2.1005307 2.1696680 2.2368937
## 2018 2.8532757 2.9254117 2.9961322 3.0654579 3.1334243
## 2019 3.7538956 3.8281287 3.9012349 3.9732072 4.0440551
```

```
## 2020 4.6837810 4.7611418 4.8375279 4.9129208 4.9873184
```

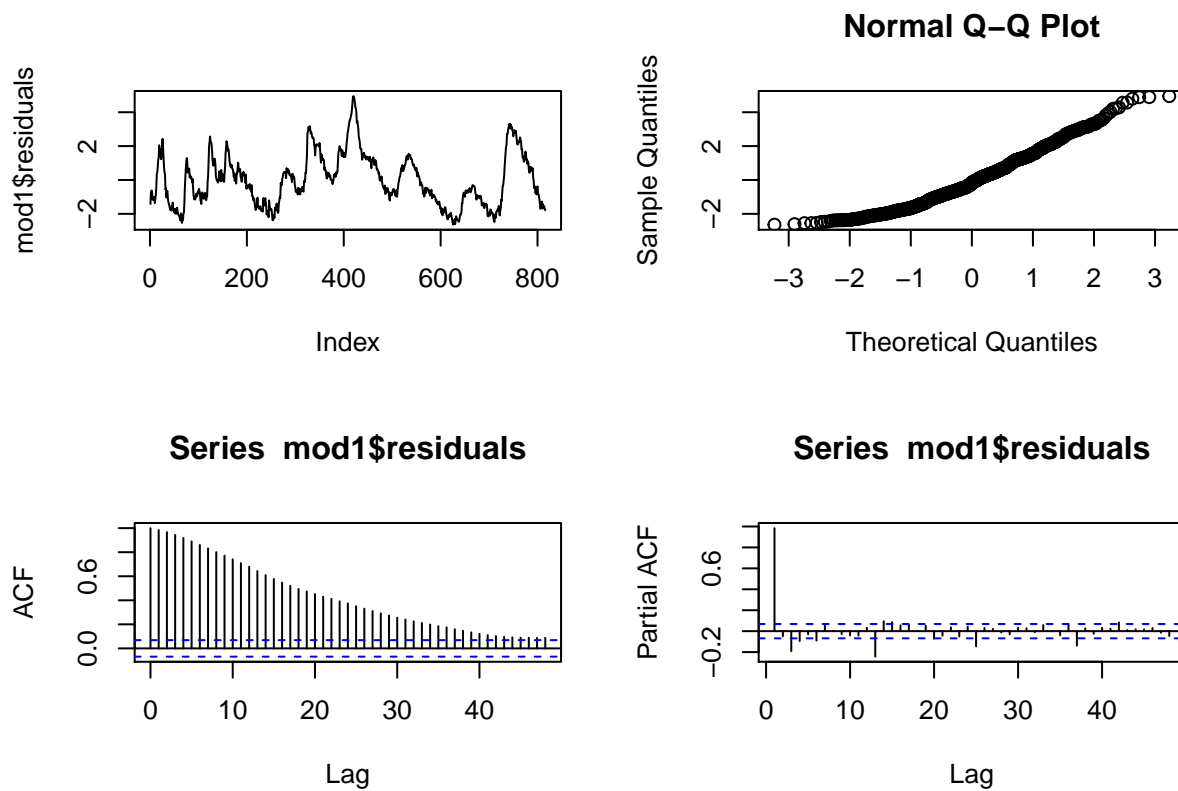
- b. Build a linear time regression and incorporate seasonal effects. Be sure to evaluate the residuals and assess this model on the bases of the assumptions of the classical linear model, and the produce a 1-year and 4-year forecast.

```
#### fit linear model w/ shorter window for data set, account for stochastic trend due to event
unemp$DATE <- as.Date(unemp$DATE)
unemp$time <- rank(unemp$DATE)
unemp$month <- months(unemp$DATE)
unemp_test <- unemp[unemp$DATE>"2015-12-01",]
unemp_train <- unemp[unemp$DATE<"2016-01-01",]
```

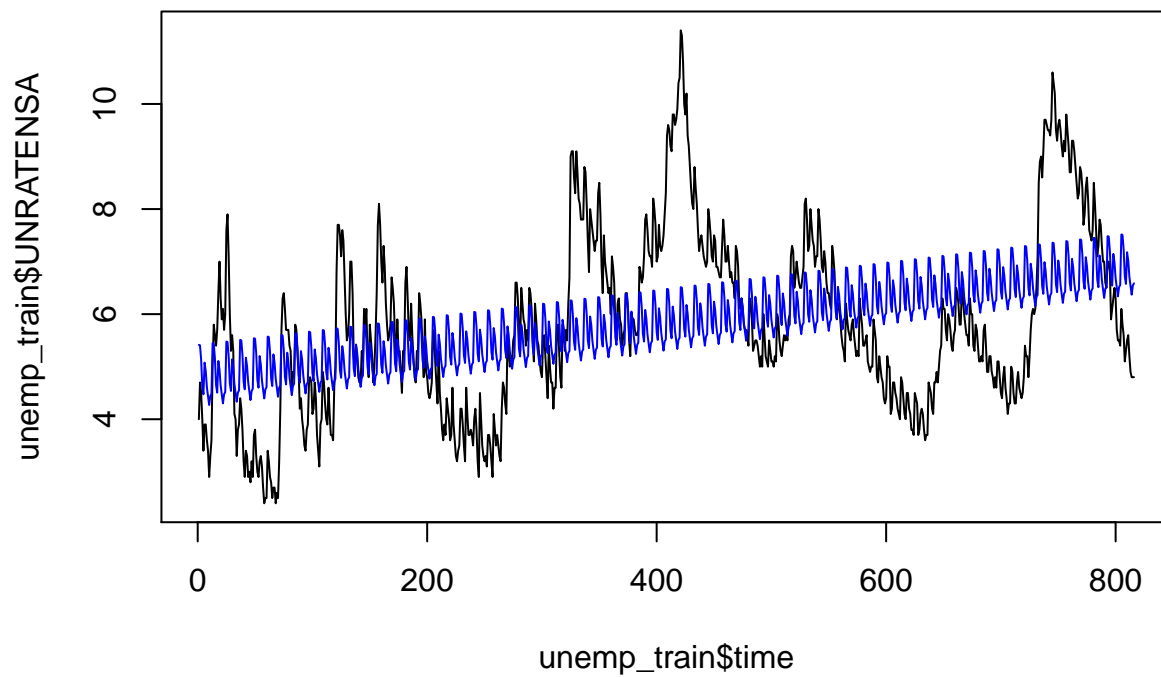
```
mod1 <- lm(UNRATENSA ~ time + month, data = unemp_train)
summary(mod1)
```

```
##
## Call:
## lm(formula = UNRATENSA ~ time + month, data = unemp_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.637 -1.156 -0.158  1.068  4.949
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   4.6084708   0.2087918  22.072 < 2e-16 ***
## time          0.0026124   0.0002291  11.403 < 2e-16 ***
## monthAugust   -0.0178027   0.2643577  -0.067  0.94633
## monthDecember -0.1547230   0.2643625  -0.585  0.55853
## monthFebruary  0.7934602   0.2643565   3.001  0.00277 **
## monthJanuary   0.8048961   0.2643570   3.045  0.00240 **
## monthJuly      0.2906921   0.2643570   1.100  0.27183
## monthJune      0.4535987   0.2643565   1.716  0.08657 .
## monthMarch     0.5143771   0.2643562   1.946  0.05203 .
## monthMay       -0.1467301   0.2643562  -0.555  0.57902
## monthNovember -0.1932870   0.2643610  -0.731  0.46490
## monthOctober  -0.3642040   0.2643597  -1.378  0.16868
## monthSeptember -0.1954151   0.2643586  -0.739  0.46000
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.541 on 803 degrees of freedom
## Multiple R-squared:  0.1833, Adjusted R-squared:  0.1711
## F-statistic: 15.02 on 12 and 803 DF,  p-value: < 2.2e-16

# check residuals - normally distributed (plot, histogram, Q-Q), white noise (ACF, PACF)
par(mfrow = c(2,2))
plot(mod1$residuals, type = "l")
qqnorm(mod1$residuals)
acf(mod1$residuals, lag.max = 48)
pacf(mod1$residuals, lag.max = 48)
```



```
# plot fitted values on training set
par(mfrow = c(1,1))
plot(unemp_train$UNRATENSA~unemp_train$time, type = "l") + lines(predict(mod1,newdata = unemp_train), col = "blue")
```




```
## numeric(0)
# forecast input data
test <- data.frame(time = 817:876, month = rep(c("January", "February", "March",
```

```

"April", "May", "June", "July",
"August", "September", "October",
"November", "December"),5))

##plot predictions
plot(predict(mod1, newdata = test[1:18,]), type="l", ylab="Unemployment Rate", ylim=c(4,8))
lines(unemp_test$UNRATENSA, col="blue")

```



```

### forecast through June 2017
accuracy(predict(mod1, newdata = test[1:18,]),unemp_test$UNRATENSA)

```

```

##           ME      RMSE      MAE      MPE      MAPE
## Test set -2.237503 2.26136 2.237503 -47.4506 47.4506

```

```

### forecast through December 2020
accuracy(predict(mod1, newdata = test[1:60,]),unemp_test$UNRATENSA)

```

```

##           ME      RMSE      MAE      MPE      MAPE
## Test set -2.237503 2.26136 2.237503 -47.4506 47.4506

```

```

mod2 <- lm(UNRATENSA ~ time + month, data = unemp_train[unemp$DATE>"2009-01-01",])
summary(mod2)

```

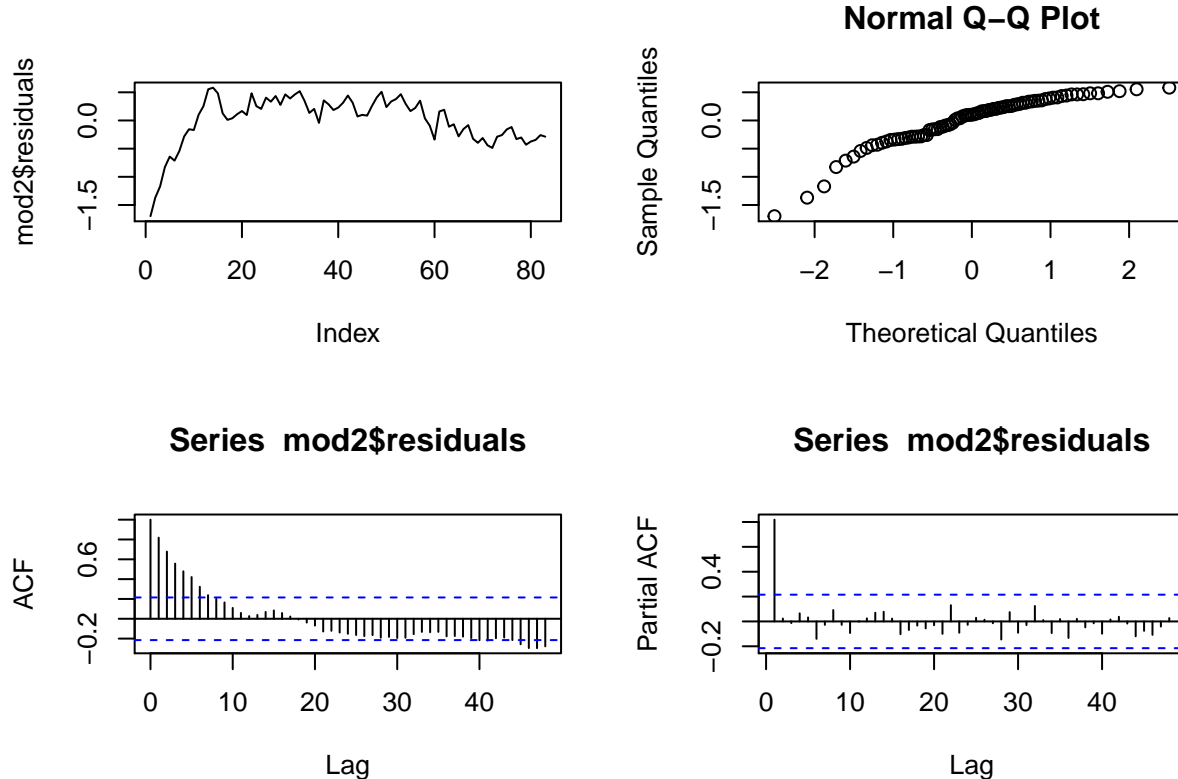
```

##
## Call:
## lm(formula = UNRATENSA ~ time + month, data = unemp_train[unemp$DATE >
##   "2009-01-01", ])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6982 -0.2830  0.1018  0.3143  0.5822
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)

```


```
## (Intercept)      55.877931      1.670875     33.442 < 2e-16 ***
## time             -0.062647      0.002152    -29.107 < 2e-16 ***
## monthAugust       0.622017      0.249220      2.496 0.01492 *
## monthDecember     0.329749      0.249666      1.321 0.19088
## monthFebruary     0.703277      0.249108      2.823 0.00619 **
## monthJanuary      1.140323      0.259322      4.397 3.83e-05 ***
## monthJuly         0.830799      0.249155      3.334 0.00137 **
## monthJune         0.696723      0.249108      2.797 0.00666 **
## monthMarch        0.537353      0.249080      2.157 0.03441 *
## monthMay          0.219790      0.249080      0.882 0.38058
## monthNovember     0.238530      0.249526      0.956 0.34240
## monthOctober      0.261598      0.249406      1.049 0.29784
## monthSeptember    0.327522      0.249304      1.314 0.19322
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.466 on 70 degrees of freedom
## (18 observations deleted due to missingness)
## Multiple R-squared:  0.9281, Adjusted R-squared:  0.9158
## F-statistic: 75.35 on 12 and 70 DF, p-value: < 2.2e-16
```

```
# check residuals - normally distributed (plot, histogram, Q-Q), white noise (ACF, PACF)
par(mfrow = c(2,2))
plot(mod2$residuals, type = "l")
qqnorm(mod2$residuals)
acf(mod2$residuals, lag.max = 48)
pacf(mod2$residuals, lag.max = 48)
```




```
# plot fitted values on training set
par(mfrow = c(1,1))
```

```
plot(unemp_train[unemp$DATE>"2009-01-01",]$UNRATENSA~unemp_train[unemp$DATE>"2009-01-01",]$time,
     type = "l") + lines(predict(mod2,newdata = unemp_train), col="blue")
```



```
## numeric(0)
# forecast input data
test <- data.frame(time = 817:876, month = rep(c("January", "February", "March",
        "April", "May", "June", "July",
        "August", "September", "October",
        "November", "December"),5))

##plot predictions
plot(predict(mod2, newdata = test[1:18,]), type="l", ylab="Unemployment Rate", ylim=c(4,8))
lines(unemp_test$UNRATENSA, col="blue")
```



```
### forecast through June 2017
accuracy(predict(mod2, newdata = test[1:18,]),unemp_test$UNRATENSA)
```

```
##           ME      RMSE      MAE      MPE      MAPE
## Test set 0.0870541 0.2283074 0.186269 2.025283 3.968705
```

```
### forecast through December 2020
accuracy(predict(mod2, newdata = test[1:60,]),unemp_test$UNRATENSA)
```

```
##           ME      RMSE      MAE      MPE      MAPE
## Test set 0.0870541 0.2283074 0.186269 2.025283 3.968705
```

- How well does your model predict the unemployment rate up until June 2017?

The linear model does not predict the unemployment rate well, with a RMSE of 2.26 on the test data set. It also significantly diverges from the observed value of 4.5% by nearly 3% points. By refitting the model to a dataset only containing unemployment rates after 2009 when the unemployment rate spiked after the financial market collapse we better estimate the new trend and our linear model predictions are more accurate.

- What does the unemployment rate look like at the end of 2020? How credible is this estimate?

The linear model fit on all data predicts an unemployment rate of 6.74% at the end of 2020. This seems fairly reasonable given previous unemployment rates but would require a significant uptick from the current trend. The model fit on data from only 2009 on predicts a rate of 1.33% which would be a historic low and does not seem credible.

- Compare this forecast to the one produced by the SARIMA model. What do you notice?

The forecast from the linear model is far less flexible and essentially estimates a flat trend with oscillating predictions based on seasonality. The linear model also is subject to the history of the data and thus if we fit with the full data set it estimates a positive trend which over the history of unemployment may be true but is certainly not reflective of the current trend.

Question 3: VAR.


You also have data on automotive car sales.

- Use a VAR model to produce a 1-year forecast on both the unemployment rate and automotive sales for 2017 in the U.S.

We are making an assumption that our predictions should be for the test set which includes 18 months of data spanning 2016 through June of 2017.

```
# split data into training and test sets
unemp.train2 = window(unemp.train, start=c(1976,1), end=c(2015,12), frequency=12)
unemp.test = window(unemp.ts, start=c(2016,1), frequency=12)
auto.train = window(auto.ts, end=c(2015,12), frequency=12)
auto.test = window(auto.ts, start=c(2016,1), frequency=12)

un_car = cbind(unemp.train2, auto.train)
plot(un_car)
```



```
# use ddlu and ddla to examine if the differenced data is stationary
ddlu.train = window(ddlu, start=c(1977,2), end=c(2015,12), frequency=12)
ddlu.test = window(ddlu, start=c(2016,1), frequency=12)
ddla.train = window(ddla, end=c(2015,12), frequency=12)
ddla.test = window(ddla, start=c(2016,1), frequency=12)
```

```
# differenced and seasonally differenced data is stationary
adf.test(ddlu.train) # p = 0.01 < 0.05, data is stationary
```

```
## Warning in adf.test(ddlu.train): p-value smaller than printed p-value
##
## Augmented Dickey-Fuller Test
##
## data: ddlu.train
## Dickey-Fuller = -6.1878, Lag order = 7, p-value = 0.01
## alternative hypothesis: stationary
```

```

pp.test(ddlu.train) # p = 0.01 < 0.05, data is stationary


## Warning in pp.test(ddlu.train): p-value smaller than printed p-value
##
## Phillips-Perron Unit Root Test
##
## data: ddlu.train
## Dickey-Fuller Z(alpha) = -615.36, Truncation lag parameter = 5,
## p-value = 0.01
## alternative hypothesis: stationary
adf.test(ddla.train) # p = 0.01 < 0.05, data is stationary

## Warning in adf.test(ddla.train): p-value smaller than printed p-value
##
## Augmented Dickey-Fuller Test
##
## data: ddla.train
## Dickey-Fuller = -11.046, Lag order = 7, p-value = 0.01
## alternative hypothesis: stationary
pp.test(ddla.train) # p = 0.01 < 0.05, data is stationary

## Warning in pp.test(ddla.train): p-value smaller than printed p-value
##
## Phillips-Perron Unit Root Test
##
## data: ddla.train
## Dickey-Fuller Z(alpha) = -518.5, Truncation lag parameter = 5,
## p-value = 0.01
## alternative hypothesis: stationary
# ccf for original data and differenced data
ccf(unemp.train2, auto.train)


```

unemp.train2 & auto.train



```
ccf(ddlu.train, ddla.train)
```

ddlu.train & ddla.train





```
# examine optimal VAR order by AIC, p = 13 shows smallest AIC  
VARselect(un_car, type='both', lag.max = 30, season=12)
```

```
## $selection
## AIC(n)  HQ(n)  SC(n) FPE(n)
##      13      4      4      13
##
## $criteria
##           1           2           3           4           5           6
## AIC(n)    5.955596    5.884010    5.792751    5.736120    5.731808    5.716316
## HQ(n)     6.063570    6.006381    5.929518    5.887284    5.897368    5.896272
## SC(n)     6.229546    6.194487    6.139754    6.119650    6.151865    6.172899
## FPE(n)  385.925893  359.272890  327.946908  309.901948  308.581674  303.853351
##           7           8           9          10          11          12
## AIC(n)    5.715299    5.712470    5.717751    5.717756    5.716798    5.645467
## HQ(n)     5.909652    5.921220    5.940897    5.955298    5.968737    5.911802
## SC(n)     6.208409    6.242107    6.283914    6.320445    6.356014    6.321210
## FPE(n)  303.562558  302.725999  304.352924  304.381877  304.121591  283.216273
##          13          14          15          16          17          18
## AIC(n)    5.613129    5.620887    5.621708    5.633582    5.640832    5.638495
## HQ(n)     5.893861    5.916015    5.931233    5.957503    5.979150    5.991209
## SC(n)     6.325398    6.369683    6.407031    6.455431    6.499208    6.533397
## FPE(n)  274.239200  276.414214  276.684368  280.037424  282.128007  281.527002
##          19          20          21          22          23          24
## AIC(n)    5.644440    5.655061    5.649709    5.663132    5.675042    5.683417
## HQ(n)     6.011551    6.036569    6.045613    6.073433    6.099739    6.122510
## SC(n)     6.575870    6.623017    6.654192    6.704142    6.752578    6.797479
## FPE(n)  283.268788  286.362313  284.907811  288.838660  292.387163  294.940892
##          25          26          27          28          29          30
## AIC(n)    5.681346    5.690357    5.702597    5.716520    5.729388    5.743720
## HQ(n)     6.134836    6.158244    6.184881    6.213199    6.240464    6.269193
## SC(n)     6.831936    6.877473    6.926240    6.976689    7.026084    7.076943
## FPE(n)  294.432166  297.206300  300.984183  305.330690  309.421621  314.034847

# season is assigned. QUESTION: should we also assign const or trend?
var1 <- VAR(un_car, p = 1, type='both', season = 12)
var13 <- VAR(un_car, p = 13, type='both', season = 12)



invisible(acf2(residuals(var1)[,1], 200))
```

Series: residuals(var1)[, 1]




```
invisible(acf2(residuals(var1)[,2], 200))
```

Series: residuals(var1)[, 2]





```
ccf(residuals(var1)[,1], residuals(var1)[,2])
```

residuals(var1)[, 1] & residuals(var1)[, 2]




```
# var13, apparently, var13 beats var1  
invisible(acf2(residuals(var13)[,1], 200))
```


Series: residuals(var13)[, 1]



```
invisible(acf2(residuals(var13)[,2], 200))
```



```
ccf(residuals(var13)[,1], residuals(var13)[,2])
```




As can be seen in the graphs above the var13 model does a far better job of producing residuals free of correlation and that simulate white noise.

```

### Forecasts
f1 <- predict(var1, n.ahead = length(unemp.test))
f13 <- predict(var13, n.ahead = length(auto.test))

# plot x.test vs forecasts, var13 shows best fitting, corresponding to residuals analysis
plot(as.vector(unemp.test), type = 'l')
lines(f1$fcst$unemp.train2[,1], col = 'red')
lines(f13$fcst$unemp.train2[,1], col = 'blue')

```




```

# forecast plots
plot(forecast(var13, 18))


```


Forecasts from VAR(13)




```
plot(forecast(var1, 18))
```

Forecasts from VAR(1)



```
# compare SARIMA and VAR
plot(as.vector(unemp.test), type = 'l')
lines(f13$fcst$unemp.train2[,1], col = 'blue')
lines(as.vector(for3$mean), col = 'red')
```



```
# check accuracy vs. SARIMA model
accuracy(f13$fcst$unemp.train2[,1], unemp.test) # RMSE = 0.285
```

```
##           ME      RMSE      MAE      MPE      MAPE      ACF1
## Test set 0.1514737 0.2848653 0.2096747 3.178804 4.343368 0.710593
##           Theil's U
## Test set 0.8722405
```

```
accuracy(for3, unemp.test) # RMSE = 0.151
```

```
##           ME      RMSE      MAE      MPE      MAPE
## Training set -0.003534511 0.2343297 0.1778710 -0.01030263 3.320626
## Test set      0.015050078 0.1506237 0.1193972 0.32696024 2.518034
##           MASE      ACF1 Theil's U
## Training set 0.2010520 -0.0004805127 NA
## Test set     0.1349576 0.3868569700 0.4453878
```

- Compare the 1-year forecast for unemployment produced by the VAR and SARIMA models, examining both the accuracy AND variance of the forecast. Do you think the addition of the automotive sales data helps? Why or why not?

Adding automotive sales data does not help the unemployment forecast, because car sales has little causal correlation with unemployment. The forecast accuracy of the VAR model on the test is less than that of the SARIMA model since RMSE increases from 0.151 to 0.285.